



Summary

Session No – 14

- A container is just a process & it has its own personal namespace
- The container looks like an operating system internally & from the OS perspective it is just a process
- One of the things the operating system gives is isolation
- Isolation gives us security & a way for organizing our resources
- The container also gives us resources (CPU, RAM, N/C) & processes
- A namespace is the one that will give isolation to the container
- Every operating system requires
 - Hardware (CPU, RAM, HD)
 - Capability to run multiple processes
 - Network i.e. IP address
 - Hostname
 - User name/login
- If we give any process a network card, CPU & RAM, username, and hostname which is also called the process tree then we can say the process is working as an operating system
- When we launch the container from the docker run command first it launches the process of the command which is given in the docker image

```
[root@ip-172-31-46-75 ~]# docker history centos:7
```

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
eeb6ee3f44bd	14 months ago	/bin/sh -c #(nop) CMD ["/bin/bash"]	0B	
<missing>	14 months ago	/bin/sh -c #(nop) LABEL org.label-schema.sc...	0B	
<missing>	14 months ago	/bin/sh -c #(nop) ADD file:b3ebbe8bd304723d4...	204MB	

```
[root@ip-172-31-46-75 ~]#
```

- Any process that is started by container technology is given a personal namespace

```
[root@ip-172-31-46-75 ~]# docker run -it --name os1 centos:7
```

```
[root@797bfc934efb /]#
```

```
[root@797bfc934efb /]#
```

```
[root@797bfc934efb /]#
```

```
[root@797bfc934efb /]#
```

```
[root@797bfc934efb /]# ps -aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	1.1	0.2	11844	2864	pts/0	Ss	16:05	0:00	/bin/bash
root	15	0.0	0.3	51748	3288	pts/0	R+	16:06	0:00	ps -aux

```
[root@797bfc934efb /]#
```

- Every process has a different hostname because on the namespace

```
[root@ip-172-31-46-75 ~]# docker attach os1
[root@797bfc934efb /]# hostname
797bfc934efb
[root@797bfc934efb /]# read escape sequence
[root@ip-172-31-46-75 ~]#
[root@ip-172-31-46-75 ~]#
[root@ip-172-31-46-75 ~]# hostname
ip-172-31-46-75.ap-south-1.compute.internal
[root@ip-172-31-46-75 ~]#
```

- As soon as we launch a container from the docker run command it will create a separate namespace for the container or process
 - Command to list namespace:- **lsns**

```
[root@ip-172-31-46-75 ~]# docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS        PORTS   NAMES
797bfc934efb   centos:7  "/bin/bash"             5 minutes ago Up 5 minutes   os1
[root@ip-172-31-46-75 ~]# lsns
NS TYPE      NPROCS  PID USER      COMMAND
4026531835  cgroup    103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531836  pid       101     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531837  user      103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531838  uts       101     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531839  ipc       101     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531840  mnt       99      1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531860  mnt       1       18 root      kdevtmpfs
4026532040  net       101     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026532213  mnt       1       2580 chrony /usr/sbin/chronyd -F 2
4026532220  mnt       2       3939 root      /bin/bash
4026532221  uts       2       3939 root      /bin/bash
4026532222  ipc       2       3939 root      /bin/bash
4026532223  pid       2       3939 root      /bin/bash
4026532225  net       2       3939 root      /bin/bash
[root@ip-172-31-46-75 ~]#
```

- As soon as we stop the container it will remove all the namespaces

```
[root@ip-172-31-46-75 ~]# docker rm -f os1
os1
[root@ip-172-31-46-75 ~]# ps -aux | grep bash
ec2-user 3769 0.0 0.4 124860 4044 pts/0    Ss   15:50   0:00 -bash
root     3832 0.0 0.4 124868 4224 pts/0    S    16:01   0:00 -bash
root     4155 0.0 0.1 119432 992 pts/0    S+   16:14   0:00 grep --color=auto bash
[root@ip-172-31-46-75 ~]# lsns
NS TYPE      NPROCS  PID USER      COMMAND
4026531835  cgroup    103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531836  pid       103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531837  user      103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531838  uts       103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531839  ipc       103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531840  mnt       101     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026531860  mnt       1       18 root      kdevtmpfs
4026532040  net       103     1 root      /usr/lib/systemd/systemd --switched-root --system --deserialize 21
4026532213  mnt       1       2580 chrony /usr/sbin/chronyd -F 2
[root@ip-172-31-46-75 ~]#
```

- **nsenter** command is used for entering the namespace
 - Command:- **nsenter -t (PID) -n**
 - -t = target
 - -n = enter network namespace

```

root@ip-172-31-46-75 ~]# nsenter -t 4204 -n
root@ip-172-31-46-75 ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.2 netmask 255.255.0.0 broadcast 172.17.255.255
    ether 02:42:ac:11:00:02 txqueuelen 0 (Ethernet)
    RX packets 13 bytes 1070 (1.0 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
root@ip-172-31-46-75 ~]#
    
```

- Linux operating system has the capability to give a network namespace & every network namespace we can give to different processes & those processes we are using with docker i.e. is the reason every docker container has a different network settings
- After we launch the container and run the ifconfig command same thing we can see the **nsenter** command also

```

root@ip-172-31-46-75 ~]# nsenter -t 4204 -n
root@ip-172-31-46-75 ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.2 netmask 255.255.0.0 broadcast 172.17.255.255
    ether 02:42:ac:11:00:02 txqueuelen 0 (Ethernet)
    RX packets 13 bytes 1070 (1.0 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
    
```

- **-a** keyword in the **nsenter** command is used to enter all the namespaces
 - Command:- **nsenter -t (PID) -a**

```

root@ip-172-31-46-75 ~]# nsenter -t 4204 -a
root@f0ae0e3e57d1 /]#
root@f0ae0e3e57d1 /]#
root@f0ae0e3e57d1 /]# hostname
f0ae0e3e57d1
root@f0ae0e3e57d1 /]# whoami
root
root@f0ae0e3e57d1 /]# ifconfig
-bash: ifconfig: command not found
root@f0ae0e3e57d1 /]#
    
```

- It means behind the scene in the docker attach command it is running **nsenter -t (PID) -a** command only
- If in one operating system we launch multiple processes all the processes by default share underline hardware resources
- Whenever we run the container whatever hardware resources we have on the base system are shared with the container
 - On the base system

- Command to list CPU:- **lscpu**

```
[root@ip-172-31-46-75 ~]# lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
CPU(s):                 1
On-line CPU(s) list:    0
Thread(s) per core:     1
Core(s) per socket:     1
Socket(s):              1
NUMA node(s):          1
Vendor ID:              GenuineIntel
CPU family:              6
```

- Command to see ram:- **free -m**

```
[root@ip-172-31-46-75 ~]# free -m
              total        used         free       shared    buff/cache       available
Mem:           964          190          147           0           627           617
Swap:           0           0            0
```

- Inside container

- Command to list CPU:- **lscpu**

```
[root@c0c666321dbd /]# lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
CPU(s):                 1
On-line CPU(s) list:    0
Thread(s) per core:     1
Core(s) per socket:     1
Socket(s):              1
NUMA node(s):          1
Vendor ID:              GenuineIntel
CPU family:              6
Model:                  63
Model name:              Intel(R) Xeon(R) CPU E5-2676 v3 @ 2.40GHz
```

- Command to see ram:- **free -m**

```
[root@ip-172-31-46-75 ~]# docker run -it centos:7
[root@c0c666321dbd /]# free -m
              total        used         free       shared    buff/cache       available
Mem:           964          204          132           0           628           603
Swap:           0           0            0
```

- Namespace does not work for hardware resources like CPU, RAM & HD
- With the help of **cgroup**, we can restrict our resources with the process

- With the help of the **memory** keyword in the run command, we can limit the memory for the docker container

```
[root@ip-172-31-46-75 ~]# docker run -it --name os5 --memory 40M centos:7
[root@7baf6ca8d72e /]#
[root@7baf6ca8d72e /]#
```

- Whenever we run the **docker run -it ubuntu14.04** command
 - It will start the process bash and give a personal namespace & we can see the namespace with the **lsns** command
 - -it means after the process start take us inside the namespace there we can use the **nsenter** command
 - Because docker has its own network namespace we can see different Ip addresses inside it
 - The container has its own / drive and the data inside is coming from the image
 - Entire hardware or resources inside the container is coming from a base operating system but we can restrict it with the help of **cgroup**
- Behind the scene, the containers is launched by the container run time program called **runc**
 - Command:- **docker info**

```
Plugins:
Volume: local
Network: bridge host ipvlan macvlan null overlay
Log: awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog
Swarm: inactive
Runtimes: io.containerd.runc.v2 io.containerd.runtime.v1.linux runc
Default Runtime: runc
Init Binary: docker-init
containerd version: 10c12954828e7c7c9b6e0ea9b0c02b01407d3ae1
runc version: 1e7bb5b773162b57333d57f612fd72e3f8612d94
init version: de40ad0
Security Options:
seccomp
Profile: default
```