



Summary

Session No – 17

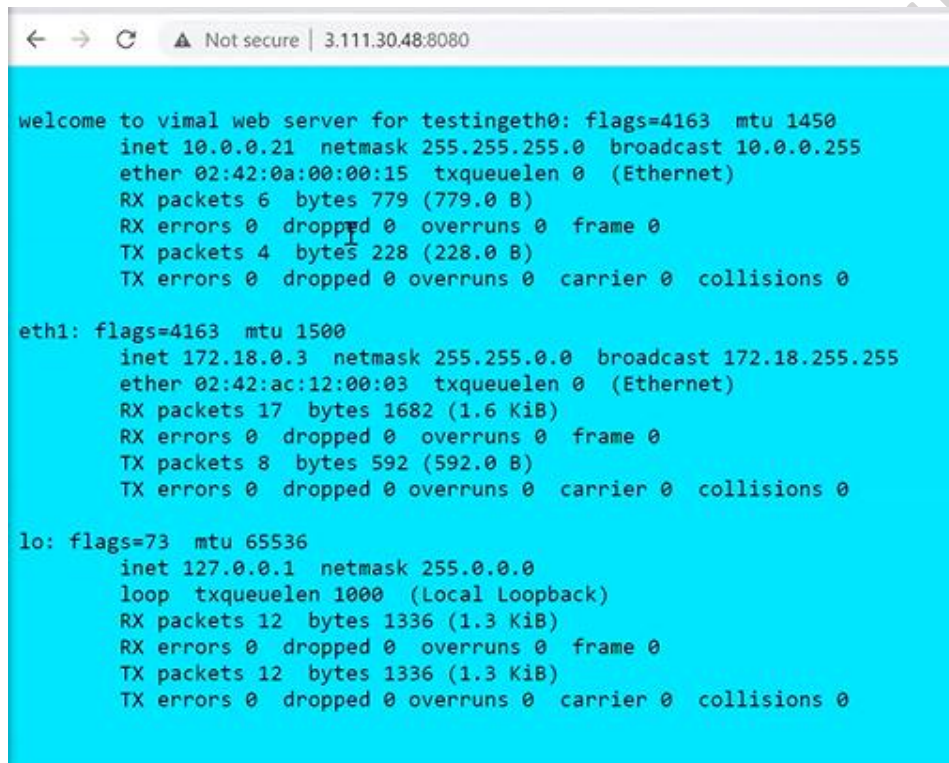
- We can launch the container with
 - Docker run command
 - Docker compose command
 - Swarm cluster
- The docker engine will give the runtime to the container with the help of a program called run c
- By default, docker gives networking in such a way that all the containers running on the same docker host can connect to each other. This kind of networking is called the bridge network
- One issue we face in networking is if we try to connect containers running on a different host, we don't have connectivity
- In multi-tier architecture, we are running our different services on the containers which are running on different host
- The overlay network & overlay driver has the capability to work as a single network on the distributed host
- As soon as we create the cluster in Swarm it automatically creates the overlay network for us & as we join more nodes automatically our distributed network will expand.
- The master node also works as a worker node
- If we list the network on the node, we can see swarm has created the overlay network

```
[root@ip-172-31-2-3 ~]# docker network ls
NETWORK ID      NAME                DRIVER              SCOPE
e7ab18ac33b4    bridge             bridge              local
64583e64fbe3    docker_gwbridge     bridge              local
444deb95af55    host               host                local
11u0hw35yy1x    ingress            overlay             swarm
7be284e23a07    none               null                local
[root@ip-172-31-2-3 ~]#
```

- Creating service on the master node

```
[root@ip-172-31-2-3 ~]# docker service ls
ID            NAME          MODE          REPLICAS    IMAGE          PORTS
[root@ip-172-31-2-3 ~]# docker service create --name myweb --publish 8080:80 vimal13/apache-webserver-php
laysztegk8tst5oelca6zzod
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
[root@ip-172-31-2-3 ~]# docker service ls
ID            NAME          MODE          REPLICAS    IMAGE          PORTS
laysztegk8t   myweb         replicated    1/1          vimal13/apache-webserver-php:latest    *:8080->80/tcp
[root@ip-172-31-2-3 ~]#
```

- The interesting thing is our container is running on the worker node & if we try to connect the website with the master's IP address, we can see the website



The screenshot shows a web browser window with the address bar displaying "Not secure | 3.111.30.48:8080". The main content area is a bright blue terminal window showing network statistics for three interfaces: eth0, eth1, and lo. The output is as follows:

```
welcome to vimal web server for testing
eth0: flags=4163 mtu 1450
    inet 10.0.0.21 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 02:42:0a:00:00:15 txqueuelen 0 (Ethernet)
    RX packets 6 bytes 779 (779.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4 bytes 228 (228.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163 mtu 1500
    inet 172.18.0.3 netmask 255.255.0.0 broadcast 172.18.255.255
    ether 02:42:ac:12:00:03 txqueuelen 0 (Ethernet)
    RX packets 17 bytes 1682 (1.6 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 8 bytes 592 (592.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

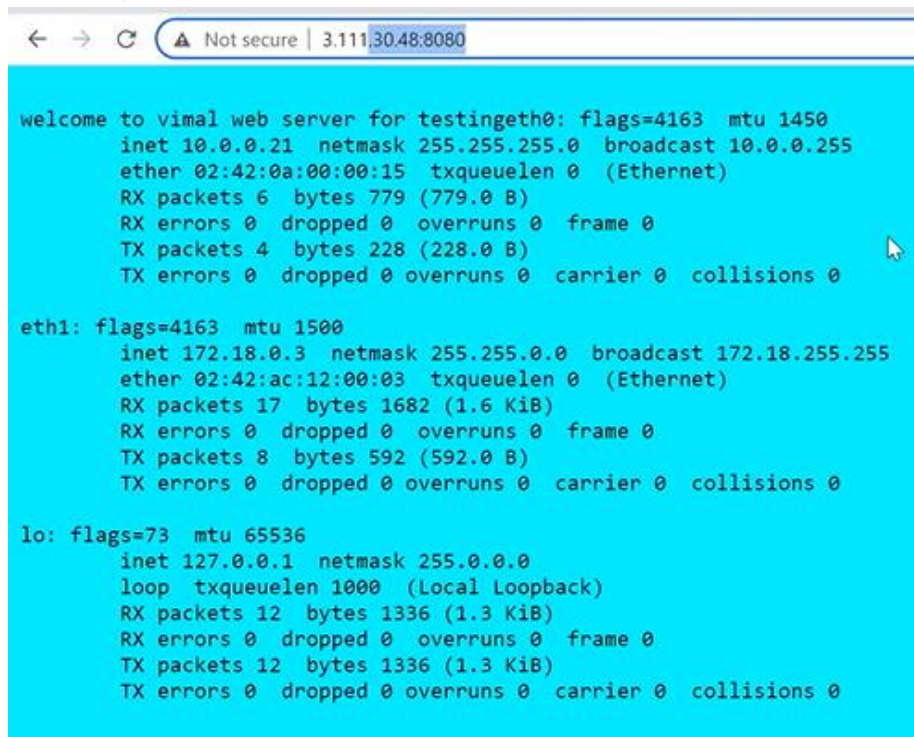
lo: flags=73 mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 12 bytes 1336 (1.3 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 12 bytes 1336 (1.3 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

- If we inspect the overlay network, we can see the network is maintaining the peers or nodes

```
{
  "com.docker.network.driver.overlay.vxlanid_list": "4096"
},
"Labels": {},
"Peers": [
  {
    "Name": "7702d1c78ef7",
    "IP": "172.31.2.3"
  },
  {
    "Name": "24fd867e8c48",
    "IP": "172.31.3.99"
  },
  {
    "Name": "12ba87d5f0f0",
    "IP": "172.31.4.217"
  }
]
}

root@ip-172-31-2-3 ~]#
```

- Even if we try to connect to the website with the worker node's IP address, we can see the website



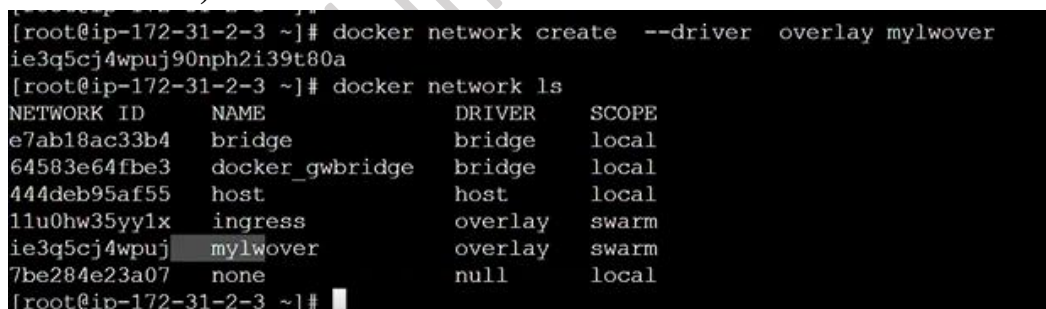
A screenshot of a web browser window. The address bar shows 'Not secure | 3.111.30.48:8080'. The main content area displays a terminal output of network configuration for a Vimal web server. The output shows details for three network interfaces: vimal0, eth1, and lo. Each interface lists its flags, MTU, IP address, netmask, broadcast address, ether address, and statistics for RX and TX packets, bytes, errors, dropped, overruns, frame, carrier, and collisions.

```
welcome to vimal web server for testing
vimal0: flags=4163 mtu 1450
inet 10.0.0.21 netmask 255.255.255.0 broadcast 10.0.0.255
ether 02:42:0a:00:00:15 txqueuelen 0 (Ethernet)
RX packets 6 bytes 779 (779.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 4 bytes 228 (228.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163 mtu 1500
inet 172.18.0.3 netmask 255.255.0.0 broadcast 172.18.255.255
ether 02:42:ac:12:00:03 txqueuelen 0 (Ethernet)
RX packets 17 bytes 1682 (1.6 KiB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 8 bytes 592 (592.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73 mtu 65536
inet 127.0.0.1 netmask 255.0.0.0
loop txqueuelen 1000 (Local Loopback)
RX packets 12 bytes 1336 (1.3 KiB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 12 bytes 1336 (1.3 KiB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

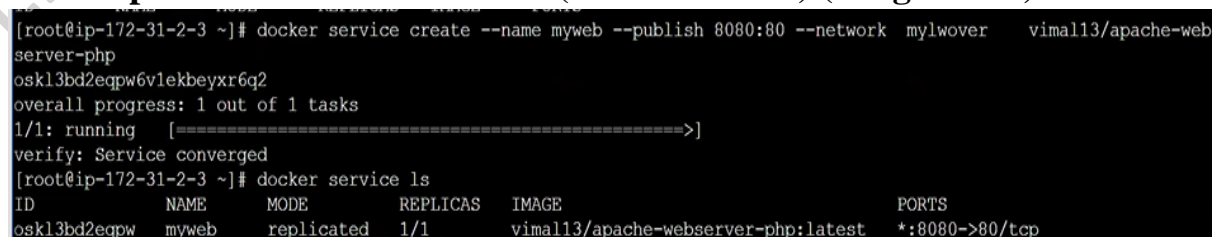
- Creating own personal overlay network with the docker network command
 - Command: - **docker network create --driver overlay (network name)**



A screenshot of a terminal window showing the creation and listing of a Docker overlay network. The first command is 'docker network create --driver overlay mylwover', which returns a long alphanumeric ID. The second command is 'docker network ls', which lists all networks with columns for NETWORK ID, NAME, DRIVER, and SCOPE. The newly created network 'mylwover' is listed with an 'overlay' driver and 'swarm' scope.

```
[root@ip-172-31-2-3 ~]# docker network create --driver overlay mylwover
ie3q5cj4wpuj90nph2i39t80a
[root@ip-172-31-2-3 ~]# docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
e7ab18ac33b4        bridge              bridge              local
64583e64f3e3        docker_gwbridge     bridge              local
444deb95af55        host                host                local
11u0hw35yy1x        ingress             overlay             swarm
ie3q5cj4wpuj        mylwover            overlay             swarm
7be284e23a07        none                null                local
[root@ip-172-31-2-3 ~]#
```

- Launching service in a custom overlay network
 - Command: - **docker service create --name (service name) --publish 8080:80 --network (network name) (image name)**



A screenshot of a terminal window showing the creation and listing of a Docker service. The first command is 'docker service create --name myweb --publish 8080:80 --network mylwover vimal13/apache-webserver-php', which shows progress and confirms the service is running. The second command is 'docker service ls', which lists all services with columns for ID, NAME, MODE, REPLICAS, IMAGE, and PORTS. The 'myweb' service is listed with a 'replicated' mode and 1/1 replicas.

```
[root@ip-172-31-2-3 ~]# docker service create --name myweb --publish 8080:80 --network mylwover vimal13/apache-webserver-php
oskl3bd2eqpw6v1ekbexr6q2
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
[root@ip-172-31-2-3 ~]# docker service ls
ID                NAME      MODE     REPLICAS  IMAGE                                  PORTS
oskl3bd2eqpw     myweb     replicated 1/1        vimal13/apache-webserver-php:latest  *:8080->80/tcp
```

- Whenever we create an overlay network or distributed network the network packet between the container is going in a clear text

- With the **encrypted** keyword in the docker network command, we can encrypt the connection

- Command :- **docker network create --opt encrypted --driver overlay (network name)**

```
[root@ip-172-31-2-3 ~]# docker network create --opt encrypted --driver overlay mysecnet
xgk0orsscer30g2mmis5thkg
[root@ip-172-31-2-3 ~]# docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
e7ab18ac33b4	bridge	bridge	local
64583e64fbe3	docker_gwbridge	bridge	local
444deb95af55	host	host	local
11u0hw35yy1x	ingress	overlay	swarm
ie3q5cj4wpuj	mylwover	overlay	swarm
xgk0orsscer	mysecnet	overlay	swarm
7be284e23a07	none	null	local

- If we inspect the network we can see the communication between the two nodes is encrypted

```

    "Gateway": "10.0.2.1"
  }
}
},
"Internal": false,
"Attachable": false,
"Ingress": false,
"ConfigFrom": {
  "Network": ""
},
"ConfigOnly": false,
"Containers": null,
"Options": {
  "com.docker.network.driver.overlay.vxlanid_list": "4098",
  "encrypted": ""
},
"Labels": null
}

```

- Creating service in a secure network

```
[root@ip-172-31-2-3 ~]# docker service create --name myweb --publish 8080:80 --network mysecnet --replicas 2 vim
al13/apache-webserver-php
ot8migvsn524hscnnc4ffnkt8
overall progress: 2 out of 2 tasks
1/2: running [=====>]
2/2: running [=====>]
verify: Service converged
[root@ip-172-31-2-3 ~]# docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
ot8migvsn524	myweb	replicated	2/2	vimal13/apache-webserver-php:latest	*:8080->80/tcp

- We have connectivity between the containers running on a different node

```
[root@7c8f4e84628d /]# ping 10.0.0.25
PING 10.0.0.25 (10.0.0.25) 56(84) bytes of data.
64 bytes from 10.0.0.25: icmp_seq=1 ttl=64 time=0.697 ms
64 bytes from 10.0.0.25: icmp_seq=2 ttl=64 time=0.626 ms
64 bytes from 10.0.0.25: icmp_seq=3 ttl=64 time=0.623 ms
64 bytes from 10.0.0.25: icmp_seq=4 ttl=64 time=0.641 ms
64 bytes from 10.0.0.25: icmp_seq=5 ttl=64 time=0.539 ms
```

- For encryption in the overlay network, it is using IPSEC protocol

- Every network card has its own physical address known as a MAC address because of which we have LAN connectivity. But we are extending the capability of Lan this concept is called VXLAN
- Docker compose has a limitation it works only on a single node
- Stack is one capability of docker with its help we can launch the container in different nodes and automate the docker-compose file
- Docker stack practical
 - Creating app.py file

```
from flask import Flask
from redis import Redis

app = Flask(__name__)
redis = Redis(host='redis', port=6379)

@app.route('/')
def hello():
    count = redis.incr('hits')
    return 'Hello World! I have been seen {} times.\n'.format(count)

if __name__ == "__main__":
    app.run(host="0.0.0.0", port=8000, debug=True)
```

- Creating requirement.txt

```
flask
redis
~
~
~
~
```

- Docker file

```
FROM python:3.4-alpine
ADD . /code
WORKDIR /code
RUN pip install -r requirements.txt
CMD ["python", "app.py"]
~
~
~
~
~
```

- Creating the image

```
[root@ip-172-31-2-3 code]# docker build -t mypy:v1 .
Sending build context to Docker daemon 4.096kB
Step 1/5 : FROM python:3.4-alpine
3.4-alpine: Pulling from library/python
8e402f1a9c57: Pulling fs layer
cda9ba2397ef: Pulling fs layer
aafecf9bbbfd: Pulling fs layer
bc2e7e266629: Waiting
e1977129b756: Waiting
```

- Docker-compose file

```
version: "3.9"

services:
  web:
    image: 127.0.0.1:5000/stackdemo
    build: .
    ports:
      - "8000:8000"
  redis:
    image: redis:alpine
```

- Launching container with docker-compose

```
[root@ip-172-31-2-3 code]# docker-compose up -d
[+] Running 7/8
  :: web Warning                                0.0s
  :: redis Pulled                               5.0s
  :: ca7dd9ec2225 Pull complete                 0.9s
  :: 83276aa4de36 Pull complete                 1.0s
  :: 731cc432e6da Pull complete                 1.3s
  :: 862de9590cc6 Pull complete                 1.9s
  :: a26b23e71d57 Pull complete                 2.0s
  :: 4b937ee5a2e0 Pull complete                 2.1s
[+] Building 0.8s (2/3)
=> [internal] load build definition from Dockerfile 0.1s
=> => transferring dockerfile: 178B                0.0s
=> [internal] load .dockerignore                   0.0s
=> => transferring context: 2B                      0.0s
=> resolve image config for docker.io/docker/dockerfile:1 0.7s
```

- Checking containers with the docker-compose ps command

```
[root@ip-172-31-2-3 code]# docker ps
```

CONTAINER ID	IMAGE	NAMES	COMMAND	CREATED	STATUS	PORTS
19af6fa72366	127.0.0.1:5000/stackdemo	code-web-1	"python app.py"	About a minute ago	Up 58 seconds	0.0.0.0:8000->8000/tcp
9c2479549103	redis:alpine	code-redis-1	"docker-entrypoint.s..."	About a minute ago	Up 58 seconds	6379/tcp

- Accessing the website with the curl command

```
[root@ip-172-31-2-3 code]# curl 127.0.0.1:8000
Hello World! I have been seen 1 times.
[root@ip-172-31-2-3 code]# curl 127.0.0.1:8000
Hello World! I have been seen 2 times.
[root@ip-172-31-2-3 code]# curl 127.0.0.1:8000
Hello World! I have been seen 3 times.
[root@ip-172-31-2-3 code]#
```

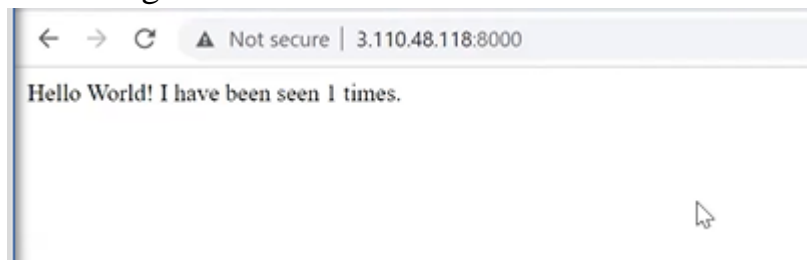
- Deploying stack

- Command:- **docker stack deploy --compose-file docker-compose.yml (Name for stack)**

```
[root@ip-172-31-2-3 code]# docker stack deploy --compose-file docker-compose.yml mypyapp
Ignoring unsupported options: build

Creating network mypyapp_default
Creating service mypyapp_web
Creating service mypyapp_redis
```

- Accessing the website from the browser



- **docker stack ls** command is used to list the stack

```
[root@ip-172-31-2-3 code]# docker stack ls
NAME          SERVICES  ORCHESTRATOR
mypyapp       2         Swarm
```

- Docker stack has the capability to launch in Kubernetes cluster also

Important link:

<https://docs.docker.com/engine/swarm/stack-deploy/>