



## Summary

### Session No- 18

- If we have a bigger workload & to manage the workload inside the containers we need multiple servers or host
- To manage multiple nodes, we need a centralized program called a manager nodes
- **docker node ls command** is used to list the nodes in the cluster
  - \* Means we are running the command from that server

```
[root@ip-172-31-2-3 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
ENGINE VERSION uvykc1f39gvckh2ycliwc2zs 20.10.17	* ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Leader
idxdku7aor668iuol8q8yq2tq 20.10.17	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active	
laxtnyozznf4g0157bpa2tldy 20.10.17	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active	

```
[root@ip-172-31-2-3 ~]#
```

- If we want to do management in a cluster it can be done only through the manager node
- Management means if we want to scale the cluster so the manager has the capability to do it
- If the client wants to connect to the web app it does not need a manager node, a client can directly connect to the worker node
- If the manager nodes go down & there are no manager nodes left, so here are manager is a single point of failure
- If we increase the manager node & any one of the manager nodes goes down we have another manager node with us this concept is called high availability
- If we don't want a single point of failure we can go for high availability

- Launching 2 new amazon Linux instances

**Launch an instance** [Info](#)

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

**Name and tags** [Info](#)

Name  
docker swarm manager 2 [Add additional tags](#)

**Application and OS Images (Amazon Machine Image)** [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

- Creating 2 instances

**Number of instances** [Info](#)

2

When launching more than 1 instance, consider [EC2 Auto Scaling](#).

**Software Image (AMI)**

Amazon Linux 2 Kernel 5.10 AMI...[read more](#)  
ami-074dc0a6f6c764218

**Virtual server type (instance type)**

t2.micro

**Firewall (security group)**

New security group

[Cancel](#) [Launch instance](#)

- Now let's say we have 3 manager nodes & anywhere any fault happens our setup can tolerate the fault this is called fault tolerance
- Creating new inbound rules for the instances

EC2 > Security Groups > sg-00e144402f1d945ed - launch-wizard-19 > Edit inbound rules

**Edit inbound rules** [Info](#)

Inbound rules control the incoming traffic that's allowed to reach the instance.

**Inbound rules** [Info](#)

Security group rule ID	Type	Protocol	Port range	Source	Description - optional
-	All traffic	All	All	Anywhere-IPv4	

[Delete](#)

- Installing docker on the instance

```
[root@ip-172-31-36-104 ~]# yum install docker -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
Resolving Dependencies
--> Running transaction check
--> Package docker.x86_64 0:20.10.17-1.amzn2.0.1 will be installed
--> Processing Dependency: runc >= 1.0.0 for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Processing Dependency: libcgrouper >= 0.40.rc1-5.15 for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Processing Dependency: containerd >= 1.3.2 for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Processing Dependency: pigz for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Running transaction check
--> Package containerd.x86_64 0:1.6.6-1.amzn2.0.2 will be installed
--> Package libcgrouper.x86_64 0:0.41-21.amzn2 will be installed
--> Package pigz.x86_64 0:2.3.4-1.amzn2.0.1 will be installed
```

- **systemctl enable docker --now** to start & enable docker permanently

```
[root@ip-172-31-36-104 ~]# systemctl enable docker --now
Created symlink from /etc/systemd/system/multi-user.target.wants/docker.service to /usr/lib/systemd/system/docker.service.
[root@ip-172-31-36-104 ~]#
```

- If we want to join a node in the existing cluster every node has different tokens, a manager has a different token & a worker node has a different token

- Creating a token for the manager node

- Command:- **docker swarm join-token manager**

```
[root@ip-172-31-2-3 ~]# docker swarm join-token manager
To add a manager to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-2w2ua5nhvc3ivhpbfnf35prhifxt3wbi96f9w9h7vrurpxr7ztf-9pdhcq3hq01vcdw4rbv4o48eg 172.31.2.3:2377

[root@ip-172-31-2-3 ~]#
```

- Running the command given by the docker swarm join-token manager on the new node

```
[root@ip-172-31-36-104 ~]# docker swarm join --token SWMTKN-1-2w2ua5nhvc3ivhpbfnf35prhifxt3wbi96f9w9h7vrurpxr7ztf-9pdhcq3hq01vcdw4rbv4o48eg 172.31.2.3:2377
This node joined a swarm as a manager.
[root@ip-172-31-36-104 ~]#
```

- Now if we run the docker info command we can see 2 manager nodes

```
Cgroup Driver: cgroupfs
Cgroup Version: 1
Plugins:
Volume: local
Network: bridge host ipvlan macvlan null overlay
Log: awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog
Swarm: active
NodeID: 2ua786v8aypnjkskwbknktw7s
Is Manager: true
ClusterID: maoobafkvdrw5oodqamdwxoe
Managers: 2
Nodes: 4
Default Address Pool: 10.0.0.0/8
SubnetSize: 24
Data Path Port: 4789
Orchestration:
Task History Retention Limit: 5
Raft:
```

- Now if we list the docker nodes we can see node is added & is reachable

```
[root@ip-172-31-2-3 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
ENGINE VERSION				
avykcl1f39gvwckh2ycliwc2zs *	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Leader
20.10.17				
idxdku7aor668iuol8q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active	
20.10.17				
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active	
20.10.17				
2ua786v8aypnjkskwobnktw7s	ip-172-31-36-104.ap-south-1.compute.internal	Ready	Active	Reachable
20.10.17				

```
[root@ip-172-31-2-3 ~]#
```

- Stopping the docker services on the leader node
  - Command:- **systemctl stop docker**

```
[root@ip-172-31-2-3 ~]# systemctl stop docker
Warning: Stopping docker.service, but it can still be activated by:
docker.socket
[root@ip-172-31-2-3 ~]#
```

- Now if we list the docker nodes we can see the other node became the leader node

```
[root@ip-172-31-36-104 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
ENGINE VERSION				
avykcl1f39gvwckh2ycliwc2zs	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Reachable
20.10.17				
idxdku7aor668iuol8q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active	
20.10.17				
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Unknown	Active	
20.10.17				
2ua786v8aypnjkskwobnktw7s *	ip-172-31-36-104.ap-south-1.compute.internal	Unknown	Active	Leader
20.10.17				

```
[root@ip-172-31-36-104 ~]#
```

- Stopping the master instance with the **init 0** command

```
[root@ip-172-31-2-3 ~]# init 0
```

- If the main master goes down and we list the node from another leader it is showing error it is a kind of cluster shutdown

```
[root@ip-172-31-36-104 ~]# docker node ls
Error response from daemon: rpc error: code = Unknown desc = The swarm does not have a leader. It's possible that too few managers are online. Make sure more than half of the managers are online.
[root@ip-172-31-36-104 ~]# docker node ls
Error response from daemon: rpc error: code = Unknown desc = The swarm does not have a leader. It's possible that too few managers are online. Make sure more than half of the managers are online.
[root@ip-172-31-36-104 ~]#
```

- The main purpose of the master node is to manage the task in the worker node. The master node maintains the current information of the worker node
- Raft is a kind of protocol or algorithm to share information between two different nodes. The main duty of the RAFT algorithm is to share the current status of the cluster with another master node
- If both masters give a different task, then there will be some kind of conflict between the masters this is called a split-brain problem
- That's the reason in the high availability where we have more than one master, we have to make one master node as a leader and others as a manager

- If the main master nodes goes down and we have two more master nodes in that case of high availability election kind of thing happens based on multiple parameters e.g. which master joins first
- $(N-1)/2$  this formula from quorum will help us to find the fault tolerance where N is no of nodes
- Now if we start the main master again and list the nodes, we can see it is reachable which means it is a passive node and the leader is an active node

```
[root@ip-172-31-2-3 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
uvykc1f39gvckh2ycliwc2zs *	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17
idxdku7aor668iuol8q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active		20.10.17
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active		20.10.17
2ua786v8aypnjkskwobnktw7s	ip-172-31-36-104.ap-south-1.compute.internal	Ready	Active	Leader	20.10.17

```
[root@ip-172-31-2-3 ~]#
```

- Installing docker on 3 rd. master node

```
[root@ip-172-31-39-233 ~]# yum install docker -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core | 3.7 kB 00:00:00
Resolving Dependencies
--> Running transaction check
--> Package docker.x86_64 0:20.10.17-1.amzn2.0.1 will be installed
--> Processing Dependency: runc >= 1.0.0 for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Processing Dependency: libcgrouper >= 0.40.rc1-5.15 for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Processing Dependency: containerd >= 1.3.2 for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Processing Dependency: pigz for package: docker-20.10.17-1.amzn2.0.1.x86_64
--> Running transaction check
--> Package containerd.x86_64 0:1.6.6-1.amzn2.0.2 will be installed
--> Package libcgrouper.x86_64 0:0.41-21.amzn2 will be installed
--> Package pigz.x86_64 0:2.3.4-1.amzn2.0.1 will be installed
```

- Joining the third master node in the cluster with token

```
[root@ip-172-31-39-233 ~]# docker swarm join --token SWMTKN-1-2w2ua5nhvc3ivhpbnf35prhifx3wbi96f9w9h7vrurpxr7ztf-9pdhcq3hq01vcdw4rbv4o48eq 172.31.2.3:2377
```

- If we list the node in a cluster we can see node is joined in a cluster and is a passive node

```
[root@ip-172-31-39-233 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
uvykc1f39gvckh2ycliwc2zs	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17
idxdku7aor668iuol8q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active		20.10.17
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active		20.10.17
2ua786v8aypnjkskwobnktw7s	ip-172-31-36-104.ap-south-1.compute.internal	Ready	Active	Leader	20.10.17
2dm7sl3wx2fsfnjz7pnxgwebv *	ip-172-31-39-233.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17

```
[root@ip-172-31-39-233 ~]#
```

- Scaling with the docker service update command
  - Command: - **docker service update --replicas=5 (service name)**

```
[root@ip-172-31-39-233 ~]# docker service update --replicas=5 mypyapp_web
```

```
mypyapp_web
overall progress: 1 out of 5 tasks
1/5: running [=====>]
2/5: No such image: 127.0.0.1:5000/stackdemo:latest
3/5: No such image: 127.0.0.1:5000/stackdemo:latest
4/5: No such image: 127.0.0.1:5000/stackdemo:latest
5/5: No such image: 127.0.0.1:5000/stackdemo:latest
```

- In the docker service update command we can change replicas of one particular service at one point in time but in the docker service scale command we can scale multiple services at one point in time



- If the leader node goes down, so according to the formula of fault tolerance we have three nodes and if one of the nodes goes down then our cluster can tolerate

```
[root@ip-172-31-39-233 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
uvykc1f39gvwckh2ycliwc2zs	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Leader	20.10.17
idxdku7aor668iuo18q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active		20.10.17
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active		20.10.17
2ua786v8aypnjkskwobnktw7s	ip-172-31-36-104.ap-south-1.compute.internal	Down	Active	Unreachable	20.10.17
2dm7sl3wx2fsfnjz7pnxgwebv *	ip-172-31-39-233.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17

- Now we can list the nodes from another active leader because our cluster has one fault tolerance

```
[root@ip-172-31-39-233 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
uvykc1f39gvwckh2ycliwc2zs	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Leader	20.10.17
idxdku7aor668iuo18q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active		20.10.17
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active		20.10.17
2ua786v8aypnjkskwobnktw7s	ip-172-31-36-104.ap-south-1.compute.internal	Down	Active	Unreachable	20.10.17
2dm7sl3wx2fsfnjz7pnxgwebv *	ip-172-31-39-233.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17

```
[root@ip-172-31-39-233 ~]# docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
pthnmxzyylfj	mypyapp_redis	replicated	1/1	redis:alpine	
3kmlw7k5lubi	mypyapp_web	replicated	5/5	127.0.0.1:5000/stackdemo:latest	*:8000->8000/tcp

- We can change any node to a master or worker node at any point in time it is called promotion or demotion
- Promoting worker node to manager node
  - Command:- **docker node promote (Hostname)**

```
[root@ip-172-31-39-233 ~]# docker node promote ip-172-31-3-99.ap-south-1.compute.internal
```

Node ip-172-31-3-99.ap-south-1.compute.internal promoted to a manager in the swarm.

```
[root@ip-172-31-39-233 ~]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
uvykc1f39gvwckh2ycliwc2zs	ip-172-31-2-3.ap-south-1.compute.internal	Ready	Active	Leader	20.10.17
idxdku7aor668iuo18q8yq2tq	ip-172-31-3-99.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17
laxtnyozznf4g0157bpa2tldy	ip-172-31-4-217.ap-south-1.compute.internal	Ready	Active		20.10.17
2ua786v8aypnjkskwobnktw7s	ip-172-31-36-104.ap-south-1.compute.internal	Down	Active	Unreachable	20.10.17
2dm7sl3wx2fsfnjz7pnxgwebv *	ip-172-31-39-233.ap-south-1.compute.internal	Ready	Active	Reachable	20.10.17

- The manager works as a master node and worker node. The good practice is to make the manager node a master-only node
- Making the manager master-only node
  - Command :- **docker node update --availability drain (hostname)**
  - Drain means if any container running on the master node stop it and launch the container on a worker node

```
2dm7sl3wx2fsfnjz7pnxgwebv ip-172-31-39-233.ap-south-1.compute.internal Ready Active Reachable
```

```
[root@ip-172-31-2-3 ~]# docker node update --availability drain ip-172-31-2-3.ap-south-1.compute.internal
```

- The drain is also used at the time of maintenance